


The hyperreduced GN kernel's cost rides on  $k$  — the mesh ( $n$ ) never enters (fixed  $m$ , rank, 10 iterations)

(a) Online solve cost vs mesh size ( $k = 8$ )



(b) Online cost vs latent dimension ( $N = 512$ )

