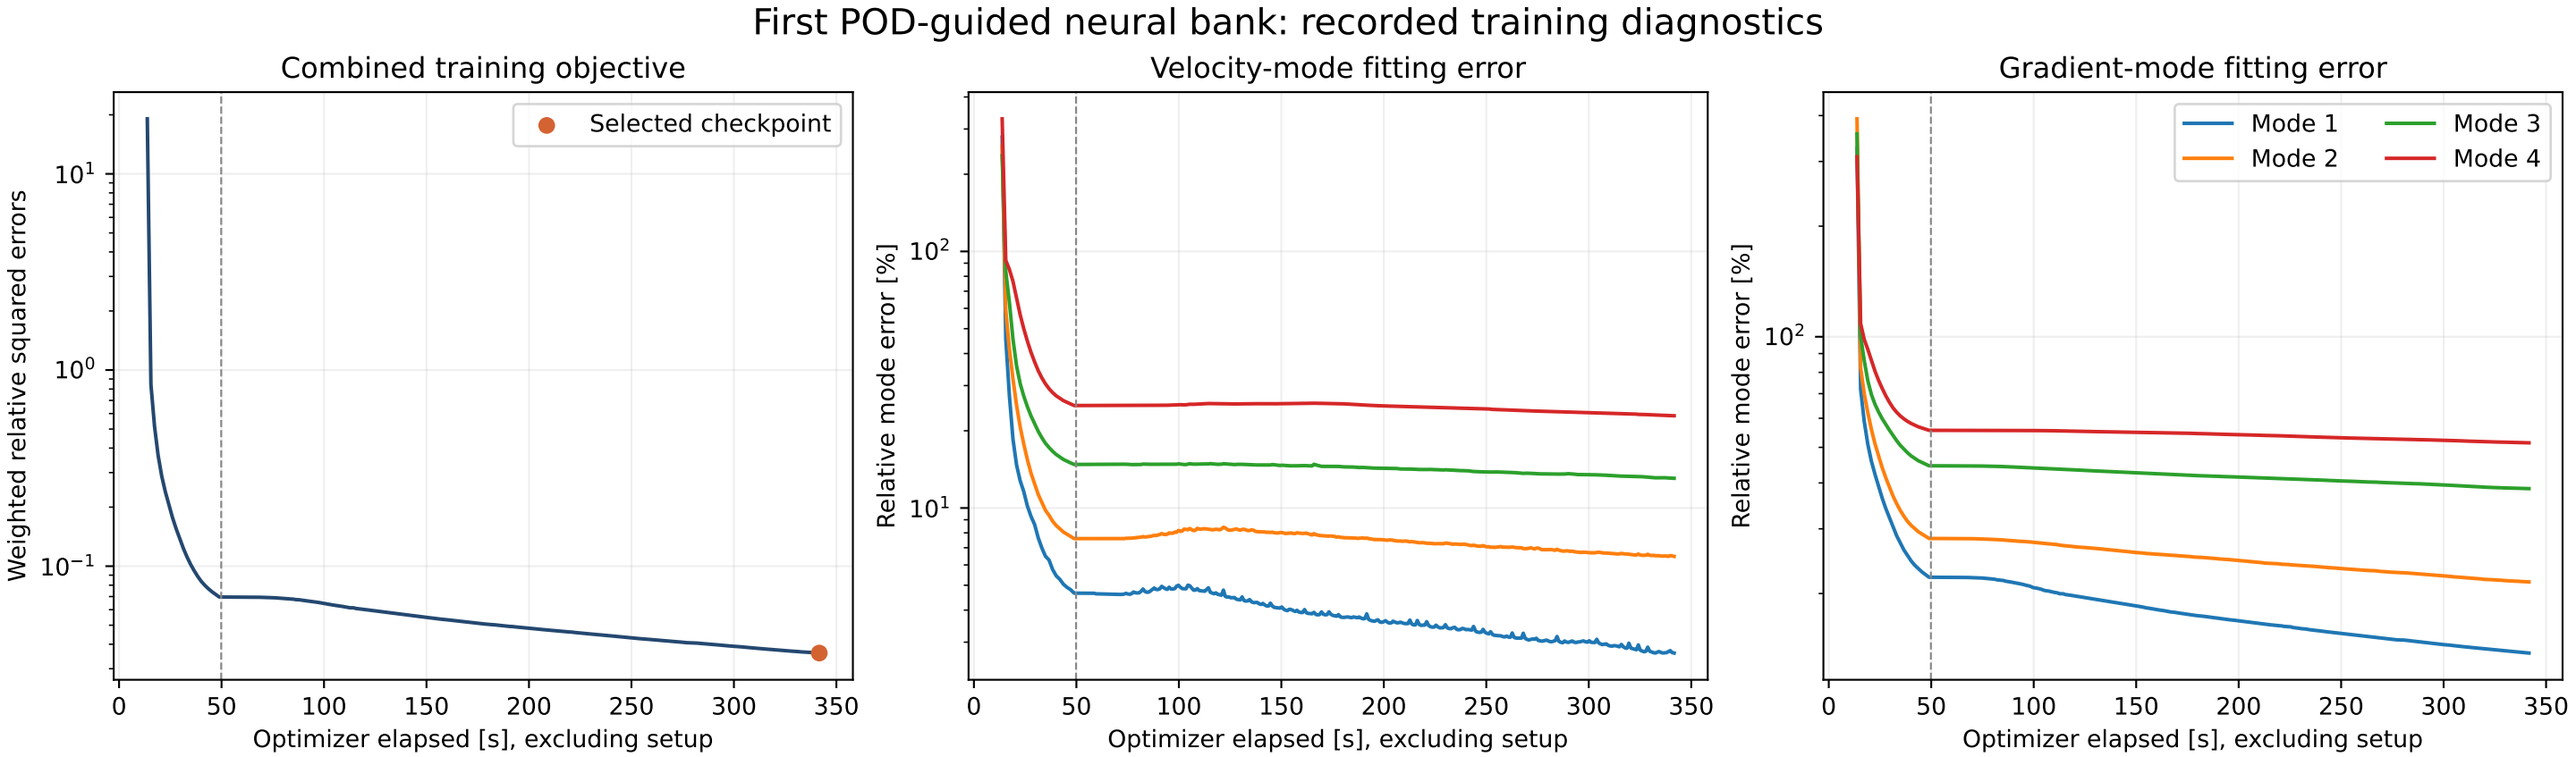
First POD-guided neural bank: recorded training diagnostics
Combined training objective
Selected checkpoint
Weighted relative squared errors
Optimizer elapsed [s], excluding setup
Velocity-mode fitting error
Relative mode error [%]
Optimizer elapsed [s], excluding setup
Gradient-mode fitting error
Mode 1
Mode 2
Mode 3
Mode 4
Relative mode error [%]
Optimizer elapsed [s], excluding setup
4,000 Adam updates and 300 accepted L-BFGS steps; dashed line marks optimizer change.
Termination: iteration_limit. These normalized training-mode errors are not actual ROM flow errors.