



A single trained decoder exposes a monotone family of offline-prepared inference-time operating points
 PROVISIONAL: the cost axis combines separate cluster jobs and may not be divided; the same-job envelope is a separate experiment

A · Burgers 2D, 256 intervals
 correction rank vs POD rank



B · Poisson 2D, 1024 intervals
 corrections eliminated exactly



C · solver-side tuning, 32 held-out cases
 DIFFERENT metric and cohort from A — do not compare by eye

